

```
In [84]: import pandas as pd
from collections import Counter
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn import metrics

from array import *
```

## Data Cleaning


```
In [85]: data = pd.read_csv('C:/Users/Frank Shi/Desktop/ADS Project 4/compas-scores-two-year')
```

```
In [86]: data.head()
```

Out[86]:

	id	name	first	last	compas_screening_date	sex	dob	age	age_cat	r
0	1	miguel hernandez	miguel	hernandez	14/08/2013	Male	18/04/1947	69	Greater than 45	O
1	3	kevon dixon	kevon	dixon	27/01/2013	Male	22/01/1982	34	25 - 45	African American
2	4	ed philo	ed	philo	14/04/2013	Male	14/05/1991	24	Less than 25	African American
3	5	marcu brown	marcu	brown	13/01/2013	Male	21/01/1993	23	Less than 25	African American
4	6	bouthy pierrelouis	bouthy	pierrelouis	26/03/2013	Male	22/01/1973	43	25 - 45	O

5 rows × 53 columns



```
In [87]: ### remove rows contains other races, update AA to be 1 and Cau to be 0
data = data[data["race"].str.contains("Other")==False]
data['race'] = data['race'].replace(['African-American', 'Caucasian'], [1, 0])
```

```
In [88]: #### update vr_charge_degree to be dummy
data['vr_charge_degree']
data['vr_charge_degree'] = data['vr_charge_degree'].fillna('0')
data['vr_charge_degree'] = data['vr_charge_degree'].str.contains(pat = '0')
```

```
In [89]: Counter(data['vr_charge_degree'])
```

Out[89]: Counter({False: 781, True: 6056})

```
In [90]: ### DROP the following columns  
### Drop the dates and columns contains too many missing values  
df = data.drop(['type_of_assessment', 'id', 'name', 'first', 'last', 'compas_score'])  
df.head()
```

```
In [93]: df.isna().sum()
```

```
Out[93]: age                                0
         race                                0
         juv_fel_count                       0
         decile_score                        0
         juv_misd_count                      0
         juv_other_count                    0
         priors_count                       0
         days_b_screening_arrest            0
         c_days_from_compas                 0
         is_recid                           0
         is_violent_recid                   0
         decile_score.1                     0
         v_decile_score                     0
         priors_count.1                     0
         start                              0
         end                                0
         event                              0
         two_year_recid                     0
         sex_Female                         0
         sex_Male                          0
         age_cat_25 - 45                    0
         age_cat_Greater than 45            0
         age_cat_Less than 25               0
         c_charge_degree_F                  0
         c_charge_degree_M                  0
         v_score_text_High                  0
         v_score_text_Low                   0
         v_score_text_Medium                0
         score_text_High                    0
         score_text_Low                     0
         score_text_Medium                  0
         dtype: int64
```

```
In [94]: ### divide the dataset into 2 dataset by races
```

```
df_cau = df[df["race"] == 0]
df_aa = df[df["race"] == 1]
print(df_cau.shape[0])
print(df_aa.shape[0])
print('Number of Cau race Commit a Crime in 2 years', df_cau[df_cau["two_year_recid"] == 1].shape[0])
print('Number of AA race Commit a Crime in 2 years', df_aa[df_aa["two_year_recid"] == 1].shape[0])
print('Percentage of Cau race Commit a Crime in 2 years', df_cau[df_cau["two_year_recid"] == 1].shape[0]/df_cau.shape[0])
print('Percentage of AA race Commit a Crime in 2 years', df_aa[df_aa["two_year_recid"] == 1].shape[0]/df_aa.shape[0])
```

```
2454
```

```
3696
```

```
Number of Cau race Commit a Crime in 2 years 966
```

```
Number of AA race Commit a Crime in 2 years 1901
```

```
Percentage of Cau race Commit a Crime in 2 years 0.39364303178484106
```

```
Percentage of AA race Commit a Crime in 2 years 0.5143398268398268
```

```
In [95]: ## drop the race column
df_cau = df_cau.drop('race', axis=1)
df_aa = df_aa.drop('race', axis=1)
```

```
In [96]: ###base model
df = df.drop('race', axis=1)
X = df.drop("two_year_recid", axis=1)
y = df["two_year_recid"]
```

```
In [97]: ## base model
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

log = LogisticRegression()
log.fit(X_train, y_train)

y_pred = log.predict(X_test)
accuracy = metrics.accuracy_score(y_test, y_pred)
accuracy
```

C:\Users\Frank Shi\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[97]: 0.9707602339181286
```

## Local Preferential Sampling method (Logitic regression)

**The data splitting framework:** We first split the data into 2 datasets by races. Within each race, we split the data into training data and testing data. We only applied local preferential sampling method on the 2 training datasets (One for African American and one for Caucasian). We combined the 2 updated training datasets, and build a new classifier based on this combined training dataset. Also, We combined the testing datasets from 2 races, and calculated overall accuracy and calibrations based on this combined testing set from 2 races.

**Race: African American.**

```
In [98]: ###data split for 2 races

### Caucasian
X_cau = df_cau.drop("two_year_recid", axis=1)
y_cau = df_cau["two_year_recid"]
X_train_cau, X_test_cau, y_train_cau, y_test_cau = train_test_split(X_cau, y_cau,

### African American
X_aa = df_aa.drop("two_year_recid", axis=1)
y_aa = df_aa["two_year_recid"]
X_train_aa, X_test_aa, y_train_aa, y_test_aa = train_test_split(X_aa, y_aa, test_
```

```
In [99]: ### Initial Logistic regression on training data for African American
log_aa = LogisticRegression()
log_aa.fit(X_train_aa, y_train_aa)

y_pred_aa = log_aa.predict(X_test_aa)
accuracy = metrics.accuracy_score(y_test_aa, y_pred_aa)
print('ACC for AA without resampling: ', accuracy) ### AA represents African Amer
```

ACC for AA without resampling: 0.9445945945945946

C:\Users\Frank Shi\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
In [100]: ## probability table construction
## dd is the table contains logits
dd = log_aa.predict_proba(X_train_aa)
(abs(dd[1,0] - dd[1,1])) ### check the difference of the second row
```

Out[100]: 0.9980437166322205

```
In [101]: ## calculate the logit differences from previous table
logit_diff = []
for i in range(len(dd)):
    logit_diff.append(abs(dd[i,0] - dd[i,1])) ### take the absolute value
```

```
In [102]: np.array(logit_diff)[np.array(logit_diff) <= 0.4]    ###max logit = 0.65
print(len(np.array(logit_diff)[np.array(logit_diff) <= 0.4])) ### number of logit

### A list contains Trues and Falses, Length of the list equals to number of rows
position = np.array(logit_diff) <= 0.4
```

151

```
In [103]: ##Gathering index for True and False
selected_rows = [] ##### index with True
not_selected_rows = [] ### index with False

for i in range(len(position)):
    if position[i] == True :
        selected_rows.append(i)
    else:
        not_selected_rows.append(i)
```

```

In [104]: ### rows with distance below threshold
selected_X = X_train_aa.iloc[selected_rows, ] ### with true in positions
selected_y = y_train_aa.iloc[selected_rows, ]

### rows with distance above threshold
unselected_X = X_train_aa.iloc[not_selected_rows, ] ### with true in positions
unselected_y = y_train_aa.iloc[not_selected_rows, ]

### merge X and y for selected and unselected
selected = pd.concat([selected_X, selected_y], axis=1)
unselected = pd.concat([unselected_X, unselected_y], axis=1)

### Only keep rows from selected that has two_year_recid == 0
### Duplicate kept rows by 4
selected = selected[selected.two_year_recid == 0] #####remain the labels with 0
repeated = pd.concat([selected]*4, ignore_index=True) #### duplicate the rows 4


### merge duplicated rows and unselected rows vertically
df_aa_train_new = pd.concat([unselected, repeated], axis=0)
(df_aa_train_new)

```

Out[104]:

	age	juv_fel_count	decile_score	juv_misd_count	juv_other_count	priors_count	days_b_scre
<b>6537</b>	27	0	6	0	0	2	
<b>5665</b>	30	0	5	0	0	6	
<b>2386</b>	41	0	4	0	0	3	
<b>4986</b>	26	0	3	0	0	2	
<b>3930</b>	24	0	3	0	0	1	
...	...	...	...	...	...	...	...
<b>363</b>	51	0	2	0	0	4	
<b>364</b>	56	0	8	0	0	24	
<b>365</b>	36	0	5	0	0	15	
<b>366</b>	23	0	5	0	0	1	
<b>367</b>	53	0	5	0	0	20	

3173 rows × 30 columns



In [105]:

```
print(df_aa_train_new.shape[0])  
print( 'Number of AA race Commit a Crime in 2 years after applied local sampling'  
print('Percentage of AA race Commit a Crime in 2 years after applied local sampli
```

3173

Number of AA race Commit a Crime in 2 years after applied local sampling 1455  
Percentage of AA race Commit a Crime in 2 years after applied local sampling 0.  
45855657106838954

In [106]:

```
print(Counter(y_train_aa))  
Counter(y_train_cau)  
770/(1184 +770)
```

Counter({1: 1514, 0: 1442})

Out[106]: 0.3940634595701126

In [107]:

```
print(Counter(y_train_aa))  
1519/(1519+1437)
```

Counter({1: 1514, 0: 1442})

Out[107]: 0.5138700947225981

### Local resampling on Causian



```

In [108]: ### M Cau
log_cau = LogisticRegression()
log_cau.fit(X_train_cau, y_train_cau)

y_pred_cau = log_cau.predict(X_test_cau)
accuracy = metrics.accuracy_score(y_test_cau, y_pred_cau)
print('Acc for Cau wihtout resampling', accuracy)

dd = log_cau.predict_proba(X_train_cau)

## calculate the Logit differences
logit_diff = []
for i in range(len(dd)):
    logit_diff.append(abs(dd[i,0] - dd[i,1])) ### take the absolute value

np.array(logit_diff)[np.array(logit_diff) <= 0.3] ###max Logit = 0.65
print(len(np.array(logit_diff)[np.array(logit_diff) <= 0.3])) ### number of Logit
position = np.array(logit_diff) <= 0.3

##Gathering index for True and False
selected_rows = [] ##Gathering index for True
not_selected_rows = [] ##Gathering index for False
for i in range(len(position)):
    if position[i] == True :
        selected_rows.append(i)
    else:
        not_selected_rows.append(i)

selected_X = X_train_cau.iloc[selected_rows, ] ### with true in positions
selected_y = y_train_cau.iloc[selected_rows, ]

unselected_X = X_train_cau.iloc[not_selected_rows, ] ### with true in positions
unselected_y = y_train_cau.iloc[not_selected_rows, ]

selected = pd.concat([selected_X, selected_y], axis=1)
unselected = pd.concat([unselected_X, unselected_y], axis=1)

### Only keep rows from selected that has two_year_recid == 1
### Duplicate kept rows by c
selected = selected[selected.two_year_recid == 1] #####remain the labels with 1
repeated = pd.concat([selected]*5, ignore_index=True)

### merge duplicated rows and unselected rows vertically
df_cau_train_new = pd.concat([unselected, repeated], axis=0)
(df_cau_train_new)

```

Acc for Cau wihtout resampling 0.9592668024439919  
46

C:\Users\Frank Shi\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

Out[108]:

	age	juv_fel_count	decile_score	juv_misd_count	juv_other_count	priors_count	days_b_scre
3864	23	0	2	0	0	0	
2143	37	0	9	0	0	11	
2923	45	0	2	1	0	7	
5284	58	0	3	0	0	10	
4021	49	0	1	0	0	0	
...	...	...	...	...	...	...	...
110	47	0	5	0	0	4	
111	53	0	5	0	0	6	
112	44	0	6	0	0	8	
113	24	0	6	0	0	1	
114	32	0	6	0	0	3	

2032 rows × 30 columns

```
In [109]: print(df_cau_train_new.shape[0])
print( 'Number of CAU race Commit a Crime in 2 years after applied local sampling
print('Percentage of CAU race Commit a Crime in 2 years after applied local sampl
```

2032

Number of CAU race Commit a Crime in 2 years after applied local sampling 858

Percentage of CAU race Commit a Crime in 2 years after applied local sampling

0.422244094488189

**Local preferential massaing: Calculate the Overall ACC and Calibration**

```
In [110]: ### merge the new training sets
df_train_new_total = pd.concat([df_aa_train_new, df_cau_train_new], axis=0)
df_train_new_x = df_train_new_total.drop("two_year_recid", axis=1)
df_train_new_y = df_train_new_total["two_year_recid"]

model2 = LogisticRegression()
model2.fit(df_train_new_x, df_train_new_y)
```

C:\Users\Frank Shi\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

Out[110]: LogisticRegression()

```
In [111]: ### acc for AA
y_pred_aa = model2.predict(X_test_aa)
accuracy_aa = metrics.accuracy_score(y_test_aa, y_pred_aa)
print('Accuracy for African American:' + str(accuracy_aa))

### acc for Cau
y_pred_cau = model2.predict(X_test_cau)
accuracy_cau = metrics.accuracy_score(y_test_cau, y_pred_cau)
print('Accuracy for Cauasin:' + str(accuracy_cau))

### overall acc
print('Accuracy total:' + str((accuracy_cau+accuracy_aa)/2))

### difference (calibrition score)
print('Difference (calibrition score):' + str(abs(accuracy_cau -accuracy_aa)))
```

Accuracy for African American:0.9675675675675676  
Accuracy for Cauasin:0.9653767820773931  
Accuracy total:0.9664721748224803  
Difference (calibrition score):0.002190785490174485

## Local Massaging (logistic)

### local message for African American (AA)

In [112]: Counter(y\_train\_aa)

Out[112]: Counter({1: 1514, 0: 1442})

In [113]: print('% of AA race Commit a Crime in 2 years before apply local massage', Counter

% of AA race Commit a Crime in 2 years before apply local massage 0.5121786197564276

In [114]: *### Method 1: Local massage for African American*

*## table contains 2 logits per row*

table\_aa = log\_aa.predict\_proba(X\_train\_aa)

*###calculate the abs difference between 2 logits from above table*

logit\_diff\_aa = []

for i in range(len(table\_aa)):

logit\_diff\_aa.append(abs(table\_aa[i,0] - table\_aa[i,1]))

print('Number of obervations below threshold', len(np.array(logit\_diff\_aa)[np.array

*#### a list contains trues and falses*

position\_aa = np.array(logit\_diff\_aa) <= 0.6

*##Label update:if the index corresopding to true, we update the lable to 0*

for i in range(len(position\_aa)):

if position\_aa[i] == True :

y\_train\_aa.iloc[i] =0

print(X\_train\_aa.shape)

print(len(y\_train\_aa))

Number of obervations below threshold 253

(2956, 29)

2956

In [115]: Counter(y\_train\_aa) *##2956*

Out[115]: Counter({1: 1405, 0: 1551})

In [116]: print('Number of AA race Commit a Crime in 2 years after applied local massage',

Number of AA race Commit a Crime in 2 years after applied local massage 0.4753044654939107

## local massage for Cau

```

In [117]: ### Method 2 on Cau
y_pred_cau = log_cau.predict(X_test_cau)
table_cau = log_cau.predict_proba(X_train_cau)

## calculate the Logit differences
logit_diff_cau = []
for i in range(len(table_cau)):
    logit_diff_cau.append(abs(table_cau[i,0] - table_cau[i,1]))

print('Number of observations below threshold', len(np.array(logit_diff_cau)[np.ar
position_cau = np.array(logit_diff_cau) <= 0.5

##Label update
for i in range(len(position_cau)):
    if position_cau[i] == True :
        y_train_cau.iloc[i] =1

print(X_train_cau.shape)
print(len(y_train_cau))

```

Number of observations below threshold 83  
(1963, 29)  
1963

```

In [118]: Counter(y_train_cau) ##1960

```

```

Out[118]: Counter({0: 1152, 1: 811})

```

```

In [119]: print('Number of Cau race Commit a Crime in 2 years after applied local message',

```

Number of Cau race Commit a Crime in 2 years after applied local message 0.4131  
43148242486

```

In [120]: pd.concat([y_train_aa, y_train_cau], axis=0)

```

```

Out[120]: 6537    1
5665    1
2386    1
4986    1
3930    0
..
4516    0
7106    1
5564    1
3713    1
6607    1
Name: two_year_recid, Length: 4919, dtype: int64

```

## Overall acc and calibration

```
In [121]: ### merge the new training sets
X_total_new = pd.concat([X_train_aa, X_train_cau], axis=0)
y_total_new = pd.concat([y_train_aa, y_train_cau], axis=0)

model3 = LogisticRegression()
model3.fit(X_total_new, y_total_new)
```

C:\Users\Frank Shi\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
 n\_iter\_i = \_check\_optimize\_result(

Out[121]: LogisticRegression()

```
In [123]: ### acc for AA
y_pred_aa = model3.predict(X_test_aa)
accuracy_aa = metrics.accuracy_score(y_test_aa, y_pred_aa)
print('Accuracy for African American:' + str(accuracy_aa))

### acc for Cau
y_pred_cau = model3.predict(X_test_cau)
accuracy_cau = metrics.accuracy_score(y_test_cau, y_pred_cau)
print('Accuracy for Cauasin:' + str(accuracy_cau))

### overall acc
print('Accuracy total:' + str((accuracy_cau+accuracy_aa)/2))

### acc difference/calibration
print('Differece/Calibration:' + str(abs(accuracy_cau-accuracy_aa)))
```

Accuracy for African American:0.9297297297297298  
 Accuracy for Cauasin:0.9226069246435845  
 Accuracy total:0.9261683271866572  
 Differece/Calibration:0.007122805086145267